

EXPERIMENT NO 05

Design the architecture and implement the autoencoder model for Image Compression.

Code and output:

```

48 import torch
import torchvision
from torch import nn
import torch.nn.functional as F
import matplotlib.pyplot as plt
import numpy as np
rng = np.random.default_rng(123456)

15 [2] data = torchvision.datasets.MNIST(root='./data', download=True)
data = data.data
data = data.float() / 255.
data = data.view(-1, 1, 28, 28)
print(data.shape)

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz to /root/data/MNIST/raw/train-images-idx3-ubyte.gz
100%|#####| 9912422/9912422 [00:00<00:00, 214073853.54it/s]Extracting /root/data/MNIST/raw/train-images-idx3-ubyte.gz to /root/data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz to /root/data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|#####| 28881/28881 [00:00<00:00, 11598591.90it/s]
Extracting /root/data/MNIST/raw/train-labels-idx1-ubyte.gz to /root/data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz to /root/data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|#####| 1648877/1648877 [00:00<00:00, 79304314.98it/s]Extracting /root/data/MNIST/raw/t10k-images-idx3-ubyte.gz to /root/data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz to /root/data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|#####| 4542/4542 [00:00<00:00, 13482327.51it/s]
Extracting /root/data/MNIST/raw/t10k-labels-idx1-ubyte.gz to /root/data/MNIST/raw

torch.Size([60000, 1, 28, 28])

05 [3] class AutoEncoder(nn.Module):
    def __init__(self):
        super().__init__()
        self.encoder = nn.Sequential(
            nn.Flatten(),
            nn.Linear(28*28, 100),
            nn.ReLU(),
            nn.Linear(100, 10),
            nn.ReLU(),
        )
        self.decoder = nn.Sequential(
            nn.Linear(10, 100),
            nn.ReLU(),
            nn.Linear(100, 28*28),
            nn.Sigmoid()
        )

        def encode(self, x):
            return self.encoder(x)

        def decode(self, x):
            x = self.decoder(x)
            return x.view(-1,1,28,28)

        def forward(self, x):
            return self.decode(self.encode(x))

15 [4] model = AutoEncoder().cuda()
opt = torch.optim.Adam(model.parameters())

1m [5] for epoch in range(25):
    print(f'Epoch {epoch+1}/25')
```

```

✓ 1m [5] for epoch in range(25):
        print(f'Epoch (epoch+1)/25')
        for i in range(0, data.shape[0], 32):
            x = data[i:i+32].cuda()
            x_rec = model(x)
            loss = F.binary_cross_entropy(x_rec, x)

            opt.zero_grad()
            loss.backward()
            opt.step()

        data = data[rng.permutation(len(data))]
        print(f'\tloss: {loss.item():.4f}')

```

```

Epoch 1/25
    loss: 0.1784
Epoch 2/25
    loss: 0.1492
Epoch 3/25
    loss: 0.1418
Epoch 4/25
    loss: 0.1496
Epoch 5/25
    loss: 0.1422
Epoch 6/25
    loss: 0.1400
Epoch 7/25
    loss: 0.1547
Epoch 8/25
    loss: 0.1331
Epoch 9/25
    loss: 0.1508
Epoch 10/25
    loss: 0.1287
Epoch 11/25
    loss: 0.1538

```


```

✓ 0s [6] plt.figure(figsize=(5,10))

        for i in range(5):
            plt.subplot(5, 2, i*2+1, title=f'Train image')
            plt.imshow(np.squeeze(x[i].cpu()), cmap='gray')
            plt.axis('off')

            plt.subplot(5, 2, i*2+2, title='Reconstruction')
            with torch.no_grad(): plt.imshow(np.squeeze(x_rec[i].cpu()), cmap='gray')
            plt.axis('off')

```



```

✓ 1s [7] # Sample two random images and encode
        f = model.encode(x[0:2])
        f1, f2 = f[0].unsqueeze(0), f[1].unsqueeze(0)

        # Show reconstructions of interpolated codes
        plt.figure(figsize=(20,5))
        reconstructions = []
        for i in range(20):
            v = i/19.
            f_interp = f1*(1-v) + f2*v
            with torch.no_grad():
                x_rec_interp = np.squeeze(model.decode(f_interp).cpu())
                reconstructions.append(x_rec_interp)

        plt.subplot(2,10,1+1)
        plt.imshow(x_rec_interp, cmap='gray')
        plt.axis('off')

```

